

# Web Server in Rust

...

04/25/2016

# Team Members



Arulselvan Madhavan

---



Jimmy Ly

---

# Agenda

---

- Introduction
- Why Rust?
- Basic web server design
- Initial thread pool design
- Improved thread pool design
- Web server cache
- Logging
- Demo
- Web server performance
- Final thoughts/next steps

# Why Rust?

- Rust is a systems programming language focused on three goals: safety, speed, and concurrency.
- It maintains these goals without having a garbage collector.
- Key Feature: Rust's Ownership Model
- 



# Basic Web Server Tasks



# Parallelizable Tasks



# Thread Pool Design - 1

- Spawn a new thread for each incoming request.
- Advantages:
  - No need to worry about data sharing and safe concurrency.
  - All the data required by a thread is self-contained.
- Disadvantages:
  - Not scalable.

# Thread Pool Design - 2





# Thread Pool based design - 2

- Challenge: Requests should be safely shared among the threads in the threadpool.
- Rust provides channels based concurrency support.



# Channel based Concurrency

- When a channel is created, a transmitter and receiver is created in the memory.
- There can be 1 transmitter and multiple receivers or multiple transmitters and 1 receiver.
- By default, there can be only 1 owner for any data in memory.
- Rust allows multiple owners to the receiver end by using Atomic Reference Counting.
- Each thread gets an atomic reference count to the receiver end.

# Thread Pool Design - 2

- Solved the scalability problem with design 1.
- Problem: Cannot enforce any priority over the jobs that needs to be served.

# Thread Pool Design - 2

- Solved the scalability problem with design 1.
- Problem: Cannot enforce any priority over the jobs that needs to be served.



# Thread Pool Design - 3

- Solution: Add an additional layer of indirection.



# Web Server Cache

- Motivation: File IO is a costly operation and some files are served so often( ex: 404 page)
- Cache file contents in web server for files that are smaller than some threshold (ie 50kb)
- Associative cache using rust-concurrent-hashmap library
- Keys are path to file, values are file contents

# Web Server Cache



# Web server Cache - Design 2

- Issue: Each thread needs to acquire lock to update cache
- Solution: Create cache thread such that only cache thread updates cache. Worker threads send cache request through channel



# Web Server Cache - Design 2



# Logger thread

- Issue: All threads should write to the same log file.
- Create a separate thread for logging.
- Each log file write requests is routed to the logger threads.
- Logger thread own the file and write to the file

# Logger Thread



# Experience with Programming in Rust

- Positive experience
- Steep learning curve
- Not worried about breaking old features when adding new ones
- Runtime errors (panicking) are still possible, if we choose to be careless.
- Many functions return a Result type (An enum of Ok(value) and Err(err)).
- Handle Results with match statements.

# Auto Restart of threads

- We have several threadpools in our web server architecture.
- Each thread is provided with the capability to restart on its own, in the event of a panic(Seg fault).

# Demo

```
Document Path:      /examplefiles/file40k
Document Length:    40960 bytes

Concurrency Level:   10
Time taken for tests: 14.976 seconds
Complete requests:   100000
Failed requests:     0
Keep-Alive requests: 0
Total transferred:   4097700000 bytes
HTML transferred:    4096000000 bytes
Requests per second: 6677.23 [#/sec] (mean)
Time per request:    1.498 [ms] (mean)
Time per request:    0.150 [ms] (mean, across all concurrent requests)
Transfer rate:       267200.09 [Kbytes/sec] received

Connection Times (ms)
      min    mean[+/-sd] median    max
Connect:    0      0   0.1      0      9
Processing:  0      1   1.0      1     32
Waiting:    0      1   0.9      1     22
Total:      0      1   1.0      1     32


Percentage of the requests served within a certain time (ms)
 50%      1
 66%      1
 75%      2
 80%      2
 90%      2
 95%      3
 98%      4
 99%      6
100%     32 (longest request)

jimmy@jimmy-VirtualBox:~/dev/rust-engine/playground/web_server$
```

# Web Server Performance

- Used ApacheBench tool to perform benchmarking on the different versions of the server.
- Different combinations of scheduler thread and worker thread counts were observed, as well as a comparison between this Rust web server and the Apache web server
- 10 concurrent connections making 100,000 total requests serving either a 10kb file, 40kb file, or 80kb file (not cached).

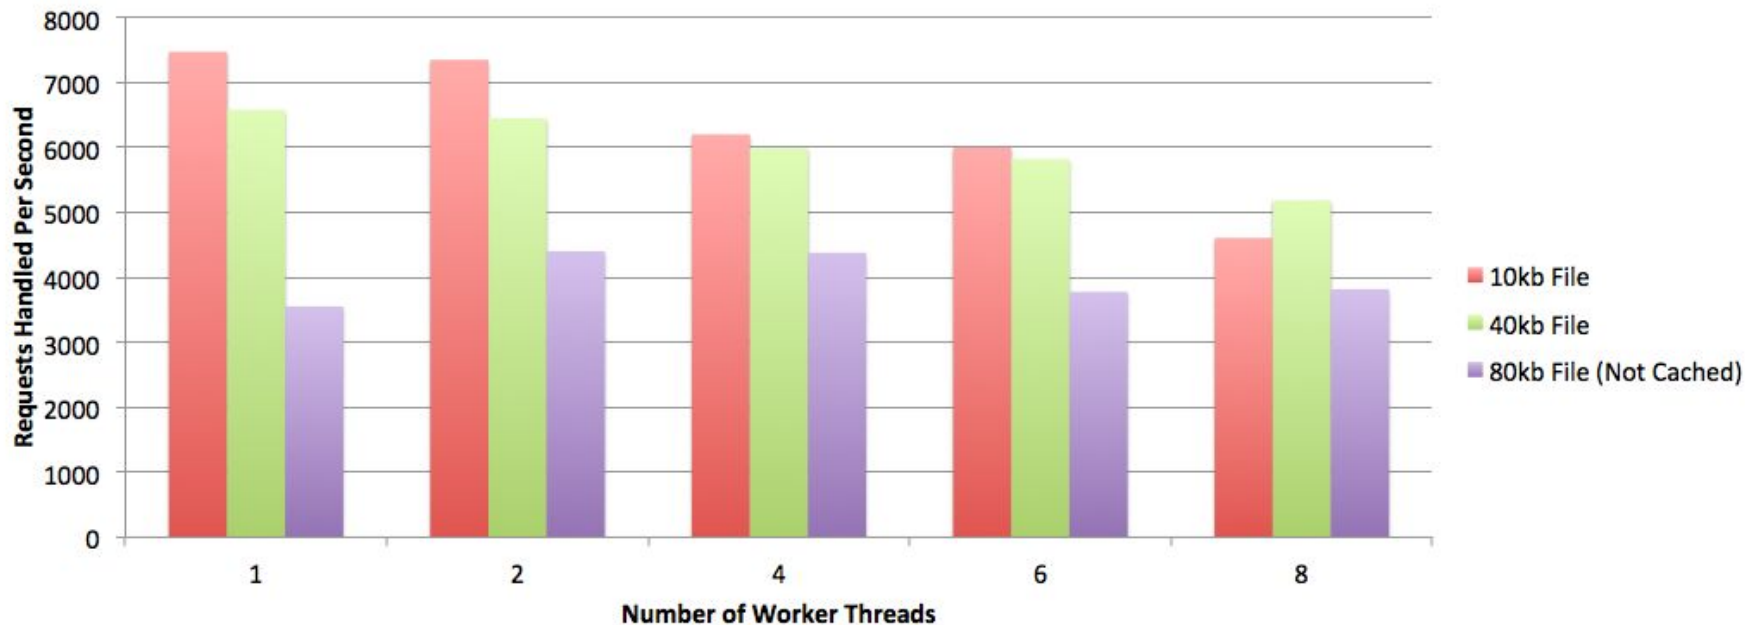
# Web Server Performance





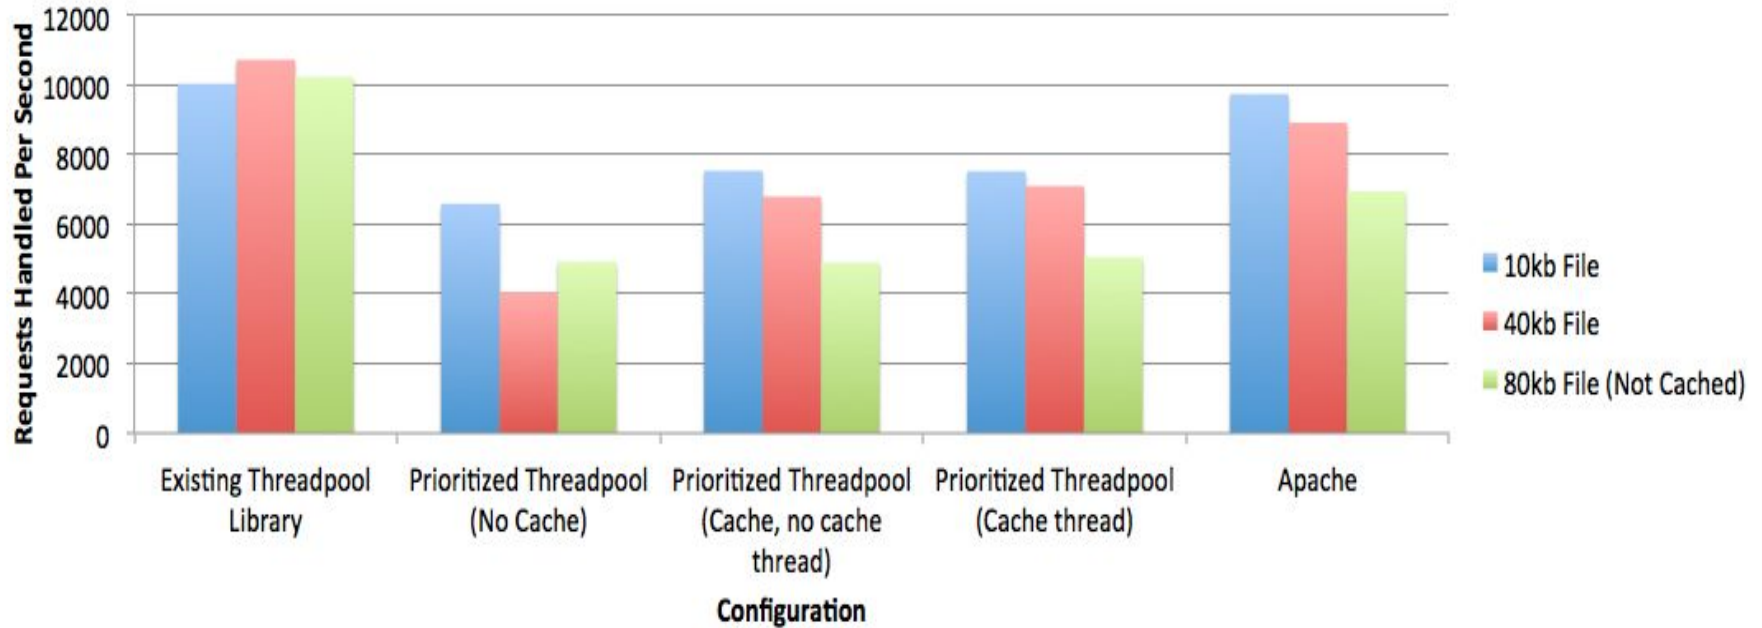
# Web Server Performance

Effect of Number of Worker Threads on Web Server Performance (Prioritized Threadpool with Cache Thread)



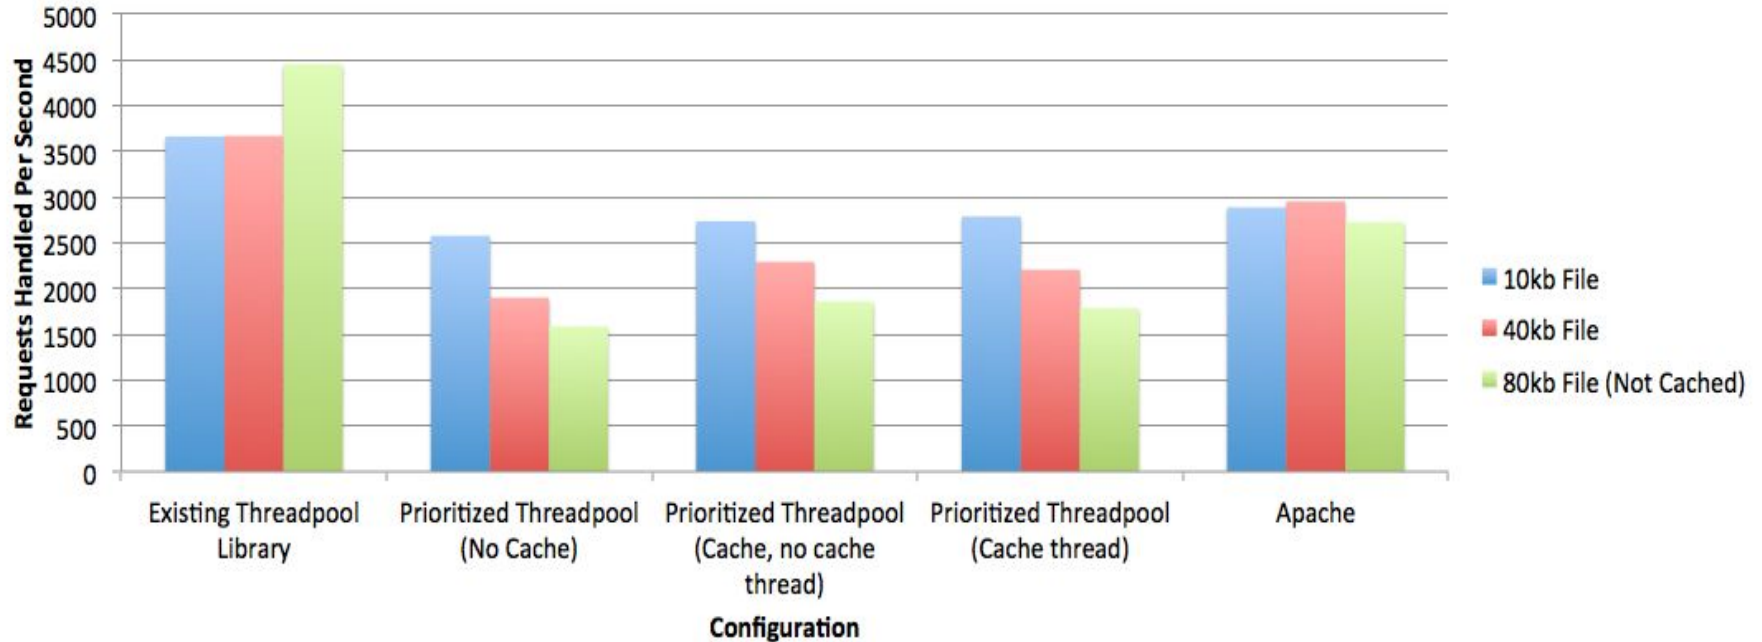
# Web Server Performance

Performance of Various Web Server Configurations and Various File Sizes



# Web Server Performance

Performance With Simultaneous Requests for Different Sized Files



# Web Server Performance

	Existing Threadpool Library	Prioritized Threadpool (No Cache)	Prioritized Threadpool (Cache, no cache thread)	Prioritized Threadpool (Cache thread)	Apache
10kb File	3659.71 requests/sec <b>9.108 sec total</b>	2575.21 requests/sec <b>12.944 sec total</b>	2734.61 requests/sec <b>12.189 sec total</b>	2785.85 requests/sec <b>11.965 sec total</b>	2885.02 requests/sec <b>11.554 sec total</b>
40kb File	3665.25 requests/sec <b>9.094 sec total</b>	1899.65 requests/sec <b>17.547 sec total</b>	2289.95 requests/sec <b>14.556 sec total</b>	2205.26 requests/sec <b>15.115 sec total</b>	2950.97 requests/sec <b>11.296 sec total</b>
80kb File (not cached)	4447.61 requests/sec <b>7.495 sec total</b>	1585.92 requests/sec <b>21.019 sec total</b>	1856.65 requests/sec <b>17.954 sec total</b>	1783.35 requests/sec <b>18.692 sec total</b>	2722.04 requests/sec <b>12.246 sec total</b>

# Final Thoughts/Next Steps

- Further tune the performance of the prioritized thread pool
- Eviction policy on cache
- Update file contents in cache if one of the files is modified
- Support for HTTP/2, IPV6, WebSockets, Directory Browsing, Virtual Hosts

**Thank You**